

ECE 637 Deep Learning Lab Exercises

Name: [Alexandre Olivé Pellicer](#)

Section 1

Exercise 1.1

1. Create two lists, `A` and `B` : `A` contains 3 arbitrary numbers and `B` contains 3 arbitrary strings.
2. Concatenate two lists into a bigger list and name that list `C` .
3. Print the first element in `C` .
4. Print the second last element in `C` via negative indexing.
5. Remove the second element of `A` from `C` .
6. Print `C` again.

```
In [ ]: # ----- YOUR CODE -----  
import numpy as np  
  
#1  
A = [1,2,3]  
B = ['Alex', 'Olive', 'Pellicer']  
  
#2  
C = A + B  
  
#3  
print(C[0])  
  
#4  
print(C[-2])  
  
#5  
rm = C.pop(1)  
  
#6  
print(C)
```

```
1  
Olive  
[1, 3, 'Alex', 'Olive', 'Pellicer']
```

Exercise 1.2

In this exercise, you will use a low-pass IIR filter to remove noise from a sine-wave signal.

You should organize your plots in a 3x1 subplot format.

1. Generate a discrete-time signal, x , by sampling a 2Hz continuous time sine wave signal with peak amplitude 1 from time 0s to 10s and at a sampling frequency of 500 Hz. Display the signal, x , from time 4s to 6s in the first row of a 3x1 subplot with the title "original signal".
2. Add Gaussian white random noise with 0 mean and standard deviation 0.1 to x and call it x_n . Display x_n from 4s to 6s on the second row of the subplot with the title "input signal".
3. Design a low-pass butterworth IIR filter of order 5 with a cut-off frequency of 4Hz, designed to filter out the noise. Hint: Use the [signal.butter](#) function and note that the frequencies are relative to the Nyquist frequency. Apply the IIR filter to x_n , and name the output y . Hint: Use [signal.filtfilt](#) function. Plot y from 4s to 6s on the third row of the subplot with the title "filtered signal".

```
In [ ]: import numpy as np                                # import the numpy packages and use a shorter
import matplotlib.pyplot as plt                          # again import the matplotlib's pyplot package
from scipy import signal                                  # import a minor package signal from scipy
plt.figure(figsize=(10, 15))                             # fix the plot size

# ----- YOUR CODE -----

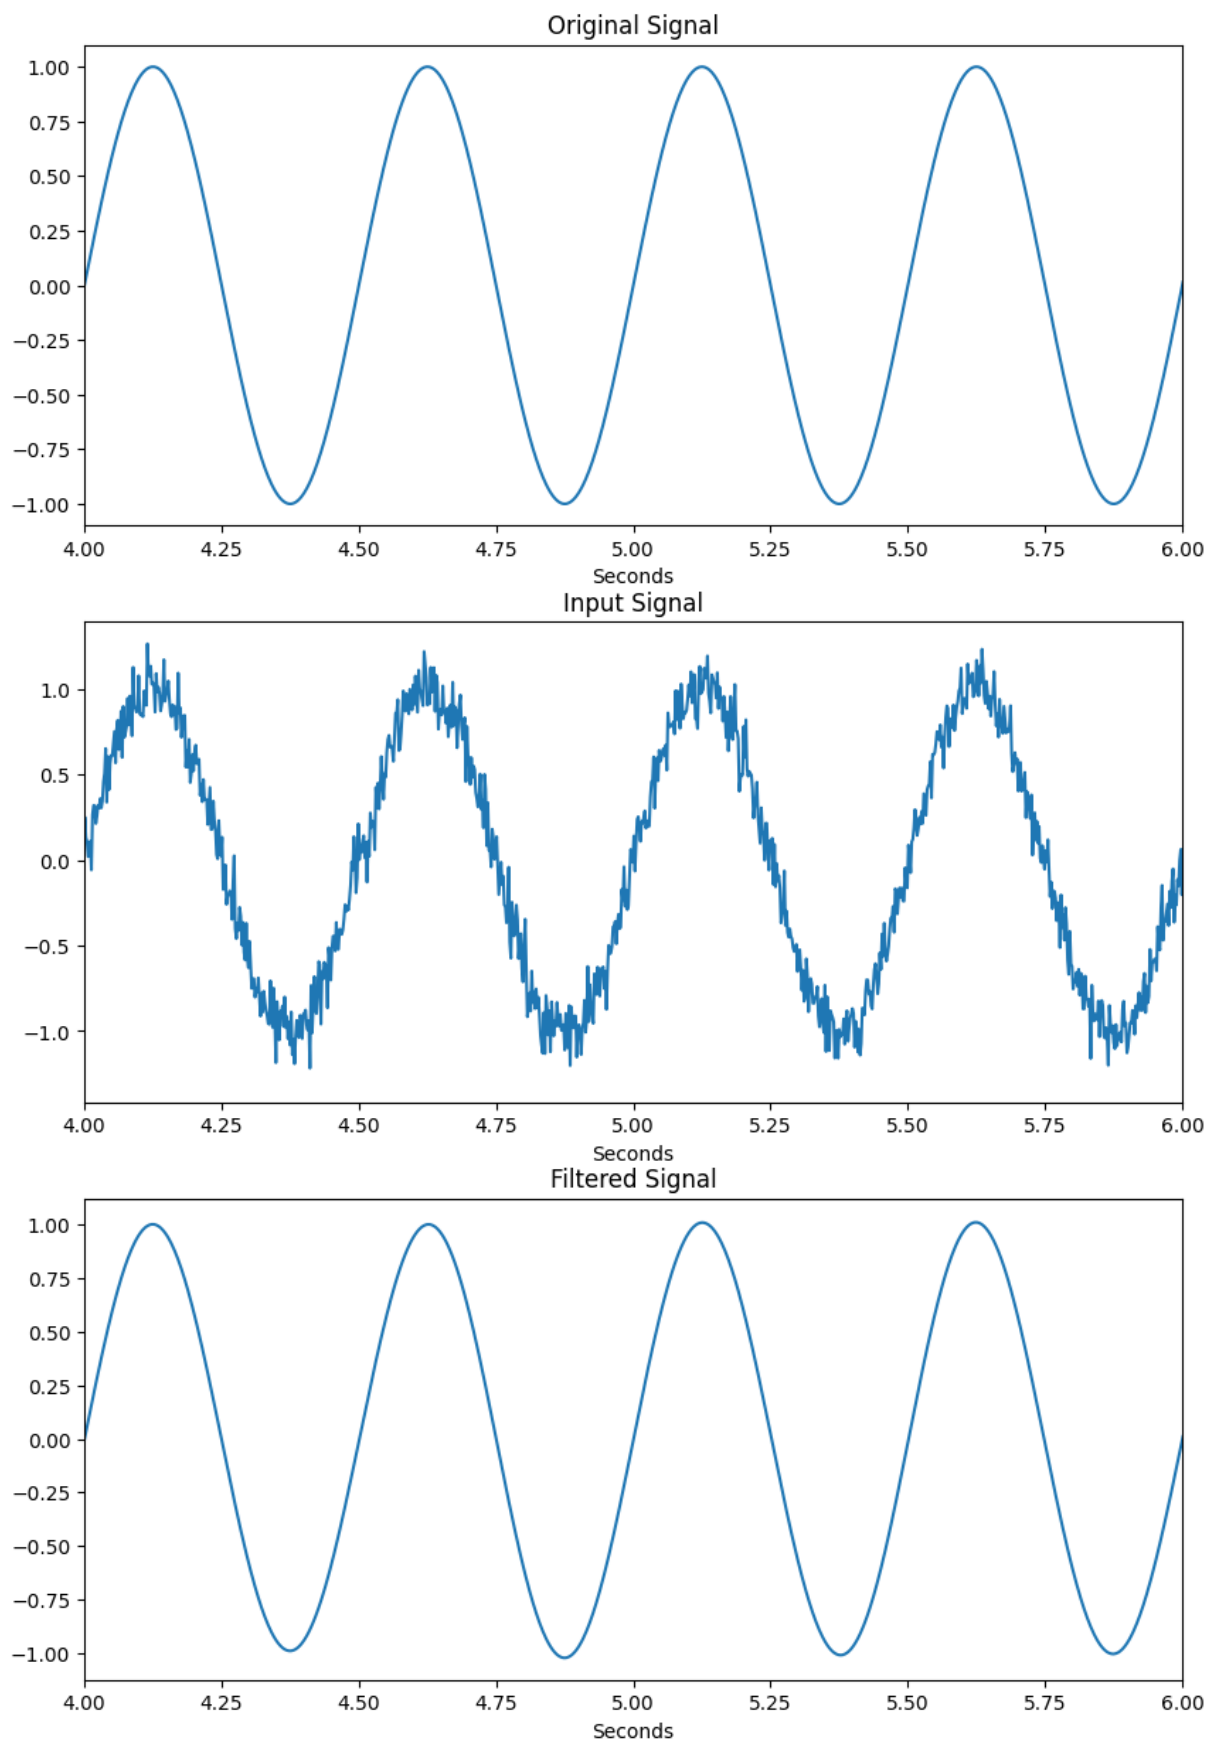
# Part 1
f = 2 # Hz
fs = 500 # sampling frequency
t_start = 0 # seconds
t_finish = 10 # seconds
num_samples = fs*(t_finish - t_start)
t = np.linspace(t_start, t_finish, num_samples)
x = np.sin(2*np.pi*f*t)
plt.subplot(3, 1, 1)
plt.plot(t, x)
plt.xlim([4, 6])
plt.title('Original Signal')
plt.xlabel('Seconds')

# Part 2
length = np.size(t);
n = np.random.randn(length)*0.1
x_n = x+n
plt.subplot(3, 1, 2)
plt.plot(t, x_n)
plt.xlim([4, 6])
plt.title('Input Signal')
plt.xlabel('Seconds')

# Part 3
cut_freq = 4
b, a = signal.butter(5, cut_freq, btype='low', fs=fs)
y = signal.filtfilt(b, a, x_n, padlen=3)
```

```
plt.subplot(3, 1, 3)
plt.plot(t, y)
plt.xlim([4, 6])
plt.title('Filtered Signal')
plt.xlabel('Seconds')

plt.show()
```



Section 2

Exercise 2.1


- Plot the third image in the test data set
- Find the corresponding label for this image and make it the title of the figure

```
In [ ]: import keras
        from keras.datasets import mnist
        (train_images, train_labels), (test_images, test_labels) = mnist.load_data()

        train_images = train_images.reshape((60000, 28, 28, 1))
        test_images = test_images.reshape((10000, 28, 28, 1))

        # ----- YOUR CODE -----
        third_image = test_images[2,:,:,:0]
        plt.imshow(third_image, cmap='gray')
        plt.title("Label: " + str(test_labels[2]))
```

```
Out[ ]: Text(0.5, 1.0, 'Label: 1')
```



Exercise 2.2

It is usually helpful to have an accuracy plot as well as a loss value plot to get an intuitive sense of how effectively the model is being trained.

- Add code to this example for plotting two graphs with the following requirements:

- Use a 1x2 subplot with the left subplot showing the loss function and right subplot showing the accuracy.
- For each graph, plot the value with respect to epochs. Clearly label the x-axis, y-axis and the title.

(Hint: The value of loss and accuracy are stored in the `hist` variable. Try to print out `hist.history` and `hist.history.keys()`.)

```
In [ ]: import keras
        from keras.datasets import mnist
        from keras import models
        from keras import layers
        from keras.utils import to_categorical

        (train_images, train_labels), (test_images, test_labels) = mnist.load_data()

        train_images = train_images.reshape((60000, 28, 28, 1))
        test_images = test_images.reshape((10000, 28, 28, 1))

        network = models.Sequential()
        network.add(layers.Flatten(input_shape=(28, 28, 1)))
        network.add(layers.Dense(512, activation='relu'))
        network.add(layers.Dense(10, activation='softmax'))

        network.summary()

        network.compile(optimizer='rmsprop', loss='categorical_crossentropy', metrics=['acc

        train_images_nor = train_images.astype('float32') / 255
        test_images_nor = test_images.astype('float32') / 255

        train_labels_cat = to_categorical(train_labels)
        test_labels_cat = to_categorical(test_labels)

        hist = network.fit(train_images_nor, train_labels_cat, epochs=5, batch_size=128)
```

Model: "sequential_27"

Layer (type)	Output Shape	Param #
flatten_13 (Flatten)	(None, 784)	0
dense_37 (Dense)	(None, 512)	401920
dense_38 (Dense)	(None, 10)	5130

=====
 Total params: 407050 (1.55 MB)
 Trainable params: 407050 (1.55 MB)
 Non-trainable params: 0 (0.00 Byte)

Epoch 1/5

469/469 [=====] - 2s 3ms/step - loss: 0.2648 - accuracy: 0.9247

Epoch 2/5

469/469 [=====] - 2s 3ms/step - loss: 0.1083 - accuracy: 0.9687

Epoch 3/5

469/469 [=====] - 2s 3ms/step - loss: 0.0712 - accuracy: 0.9784

Epoch 4/5

469/469 [=====] - 2s 3ms/step - loss: 0.0517 - accuracy: 0.9848

Epoch 5/5

469/469 [=====] - 2s 3ms/step - loss: 0.0385 - accuracy: 0.9885


```
In [ ]: import matplotlib.pyplot as plt

plt.figure(figsize=(10, 4))

# ----- YOUR CODE -----
epoch = [1, 2, 3, 4, 5]
plt.subplot(1, 2, 1)
plt.plot(epoch, hist.history['loss'])
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.title('Loss vs Epoch')

plt.subplot(1, 2, 2)
plt.plot(epoch, hist.history['accuracy'])
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.title('Accuracy vs Epoch')

plt.show()
```



Exercise 2.3

Use the dense network from Section 2 as the basis to construct a deeper network with

- 5 dense hidden layers with dimensions [512, 256, 128, 64, 32] each of which uses a ReLU non-linearity

Question: Will the accuracy on the testing data always get better if we keep making the neural network larger?

No. Not always that we make the neural network larger the accuracy will get better. There will be a moment where the neural network will be too large so that it will overfit the training data. We say that the capacity of the neural network will be too high with respect to the training data points. As a result, the model will not generalize so it will not perform well on the testing data.

```
In [ ]: import keras
        from keras import models
        from keras import layers

        # ----- YOUR CODE -----
        network = models.Sequential()
        network.add(layers.Flatten(input_shape=(28, 28, 1)))
        network.add(layers.Dense(512, activation='relu'))
        network.add(layers.Dense(256, activation='relu'))
        network.add(layers.Dense(128, activation='relu'))
        network.add(layers.Dense(64, activation='relu'))
        network.add(layers.Dense(32, activation='relu'))
        network.add(layers.Dense(10, activation='softmax'))

        network.summary()
```


Model: "sequential_28"

Layer (type)	Output Shape	Param #
=====		
flatten_14 (Flatten)	(None, 784)	0
dense_39 (Dense)	(None, 512)	401920
dense_40 (Dense)	(None, 256)	131328
dense_41 (Dense)	(None, 128)	32896
dense_42 (Dense)	(None, 64)	8256
dense_43 (Dense)	(None, 32)	2080
dense_44 (Dense)	(None, 10)	330
=====		
Total params: 576810 (2.20 MB)		
Trainable params: 576810 (2.20 MB)		
Non-trainable params: 0 (0.00 Byte)		

```
In [ ]: import keras
        from keras.datasets import mnist
        from keras.utils import to_categorical

        (train_images, train_labels), (test_images, test_labels) = mnist.load_data()

        train_images = train_images.reshape((60000, 28, 28, 1))
        test_images = test_images.reshape((10000, 28, 28, 1))

        network.compile(optimizer='rmsprop', loss='categorical_crossentropy', metrics=['acc

        train_images_nor = train_images.astype('float32') / 255
        test_images_nor = test_images.astype('float32') / 255

        train_labels_cat = to_categorical(train_labels)
        test_labels_cat = to_categorical(test_labels)

        hist = network.fit(train_images_nor, train_labels_cat, epochs=5, batch_size=128)

        test_loss, test_acc = network.evaluate(test_images_nor, test_labels_cat)
        print('test_accuracy:', test_acc)
```

```
Epoch 1/5
469/469 [=====] - 3s 4ms/step - loss: 0.3013 - accuracy: 0.
9072
Epoch 2/5
469/469 [=====] - 2s 4ms/step - loss: 0.1016 - accuracy: 0.
9686
Epoch 3/5
469/469 [=====] - 2s 4ms/step - loss: 0.0683 - accuracy: 0.
9789
Epoch 4/5
469/469 [=====] - 3s 6ms/step - loss: 0.0499 - accuracy: 0.
9845
Epoch 5/5
469/469 [=====] - 2s 4ms/step - loss: 0.0387 - accuracy: 0.
9882
313/313 [=====] - 1s 3ms/step - loss: 0.0683 - accuracy: 0.
9805
test_accuracy: 0.9804999828338623
```

Section 3

Exercise 3.1

In this exercise, you will access the relationship between the feature extraction layer and classification layer. The example above uses two sets of convolutional layers and pooling layers in the feature extraction layer and two dense layers in the classification layers. The overall performance is around 98% for both training and test dataset. In this exercise, try to create a similar CNN network with the following requirements:

- Achieve the overall accuracy higher than 99% for training and testing dataset.
- Keep the total number of parameters used in the network lower than 100,000.

```
In [ ]: import keras
        from keras import models
        from keras import layers

        network = models.Sequential()

        # ----- YOUR CODE -----
        network.add(layers.Conv2D(16, (3, 3), activation='relu', padding = 'same', input_shape=(28, 28, 1)))
        network.add(layers.MaxPooling2D((2, 2)))
        network.add(layers.Conv2D(32, (3, 3), activation='relu', padding = 'same'))
        network.add(layers.MaxPooling2D((2, 2)))
        network.add(layers.Conv2D(64, (3, 3), activation='relu', padding = 'same'))
        network.add(layers.MaxPooling2D((2, 2)))

        network.add(layers.Flatten())
        network.add(layers.Dense(128, activation='relu'))
        network.add(layers.Dense(10, activation='softmax'))

        network.summary()
```

Model: "sequential_29"

Layer (type)	Output Shape	Param #
=====		
conv2d_43 (Conv2D)	(None, 28, 28, 16)	160
max_pooling2d_19 (MaxPooling2D)	(None, 14, 14, 16)	0
conv2d_44 (Conv2D)	(None, 14, 14, 32)	4640
max_pooling2d_20 (MaxPooling2D)	(None, 7, 7, 32)	0
conv2d_45 (Conv2D)	(None, 7, 7, 64)	18496
max_pooling2d_21 (MaxPooling2D)	(None, 3, 3, 64)	0
flatten_15 (Flatten)	(None, 576)	0
dense_45 (Dense)	(None, 128)	73856
dense_46 (Dense)	(None, 10)	1290
=====		
Total params: 98442 (384.54 KB)		
Trainable params: 98442 (384.54 KB)		
Non-trainable params: 0 (0.00 Byte)		

```
In [ ]: from keras.datasets import mnist
        from keras.utils import to_categorical

        (train_images, train_labels), (test_images, test_labels) = mnist.load_data()
```

```

train_images = train_images.reshape((60000, 28, 28, 1))
train_images_nor = train_images.astype('float32') / 255
test_images = test_images.reshape((10000, 28, 28, 1))
test_images_nor = test_images.astype('float32') / 255

train_labels_cat = to_categorical(train_labels)
test_labels_cat = to_categorical(test_labels)

network.compile(optimizer='rmsprop', loss='categorical_crossentropy', metrics=['acc'])
network.fit(train_images_nor, train_labels_cat, epochs=5, batch_size=128)

test_loss, test_acc = network.evaluate(test_images_nor, test_labels_cat)
print('test_accuracy:', test_acc)

```

```

Epoch 1/5
469/469 [=====] - 3s 5ms/step - loss: 0.2746 - accuracy: 0.9137
Epoch 2/5
469/469 [=====] - 2s 5ms/step - loss: 0.0626 - accuracy: 0.9803
Epoch 3/5
469/469 [=====] - 2s 5ms/step - loss: 0.0410 - accuracy: 0.9867
Epoch 4/5
469/469 [=====] - 3s 6ms/step - loss: 0.0304 - accuracy: 0.9902
Epoch 5/5
469/469 [=====] - 2s 5ms/step - loss: 0.0239 - accuracy: 0.9924
313/313 [=====] - 1s 4ms/step - loss: 0.0263 - accuracy: 0.9902
test_accuracy: 0.9901999831199646

```

Section 4

Exercise 4.1

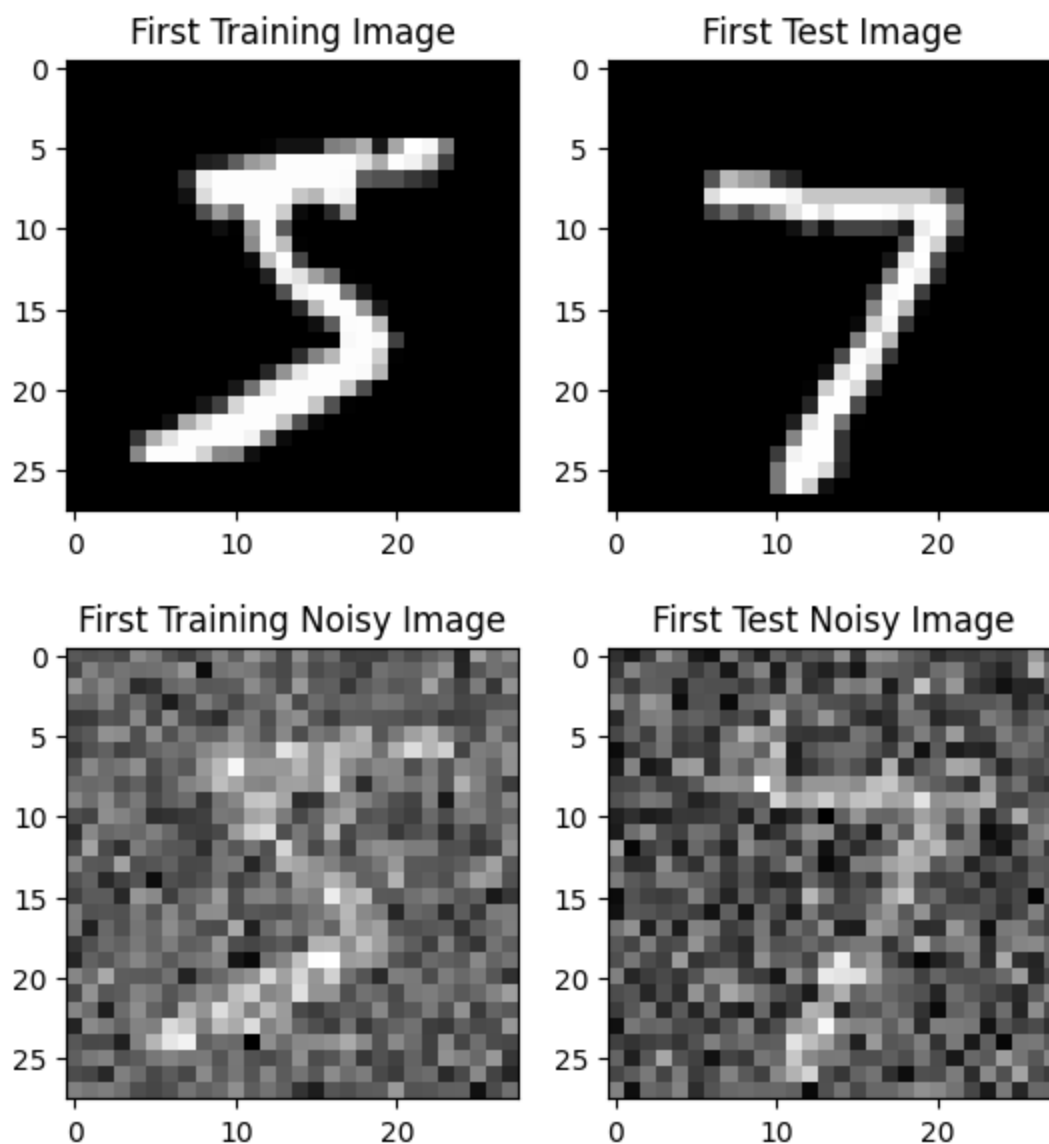
In this exercise you will need to create the entire neural network that does image denoising tasks. Try to mimic the code provided above and follow the structure as provided in the instructions below.

Task 1: Create the datasets

1. Import necessary packages
2. Load the MNIST data from Keras, and save the training dataset images as `train_images`, save the test dataset images as `test_images`
3. Add additive white gaussian noise to the train images as well as the test images and save the noisy images to `train_images_noisy` and `test_images_noisy` respectively. The noise should have mean value 0, and standard deviation 0.4. (Hint: Use `np.random.normal`)

4. Show the first image in the training dataset as well as the test dataset (plot the images in 1 x 2 subplot form)

```
In [ ]: # ----- YOUR CODE -----  
  
import keras  
from keras.datasets import mnist  
from keras import models  
from keras import layers  
from tensorflow.keras.utils import to_categorical  
import matplotlib.pyplot as plt  
import numpy as np  
  
(train_images, train_labels), (test_images, test_labels) = mnist.load_data()  
  
train_images = train_images.reshape((60000, 28, 28, 1))  
train_images_nor = train_images.astype('float32') / 255  
test_images = test_images.reshape((10000, 28, 28, 1))  
test_images_nor = test_images.astype('float32') / 255  
  
train_noise = np.random.normal(0, 0.4, train_images.shape)  
test_noise = np.random.normal(0, 0.4, test_images.shape)  
  
train_images_noisy = train_images_nor + train_noise  
test_images_noisy = test_images_nor + test_noise  
  
train_labels_cat = to_categorical(train_labels)  
test_labels_cat = to_categorical(test_labels)  
  
first_train = train_images[0,:,:,0]  
first_test = test_images[0,:,:,0]  
  
plt.subplot(1, 2, 1)  
plt.imshow(first_train, cmap='gray')  
plt.title('First Training Image')  
  
plt.subplot(1, 2, 2)  
plt.imshow(first_test, cmap='gray')  
plt.title('First Test Image')  
  
plt.show()  
  
first_noisy_train = train_images_noisy[0,:,:,0]  
first_noisy_test = test_images_noisy[0,:,:,0]  
  
plt.subplot(1, 2, 1)  
plt.imshow(first_noisy_train, cmap='gray')  
plt.title('First Training Noisy Image')  
  
plt.subplot(1, 2, 2)  
plt.imshow(first_noisy_test, cmap='gray')  
plt.title('First Test Noisy Image')  
  
plt.show()
```



Task 2: Create the neural network model

1. Create a sequential model called `encoder` with the following layers sequentially:

- convolutional layer with 32 output channels, 3x3 kernel size, and the padding convention 'same' with 'relu' activation function.
- max pooling layer with 2x2 kernel size
- convolutional layer with 16 output channels, 3x3 kernel size, and the padding convention 'same' with 'relu' activation function.
- max pooling layer with 2x2 kernel size
- convolutional layer with 8 output channels, 3x3 kernel size, and the padding convention 'same' with 'relu' activation function and name the layer as `'convOutput'`.
- flatten layer
- dense layer with output dimension as `encoding_dim` with 'relu' activation function.

2. Create a sequential model called `decoder` with the following layers sequentially:

- dense layer with the input dimension as `encoding_dim` and the output dimension as the product of the output dimensions of the `'convOutput'` layer.
- reshape layer that convert the tensor into the same shape as `'convOutput'`
- convolutional layer with 8 output channels, 3x3 kernel size, and the padding convention `'same'` with `'relu'` activation function.
- upsampling layer with 2x2 kernel size
- convolutional layer with 16 output channels, 3x3 kernel size, and the padding convention `'same'` with `'relu'` activation function.
- upsampling layer with 2x2 kernel size
- convolutional layer with 32 output channels, 3x3 kernel size, and the padding convention `'same'` with `'relu'` activation function
- convolutional layer with 1 output channels, 3x3 kernel size, and the padding convention `'same'` with `'sigmoid'` activation function

3. Create a sequential model called `autoencoder` with the following layers sequentially:

- `encoder` model
- `decoder` model

```
In [ ]: # ----- YOUR CODE -----
encoding_dim = 32

# 1
encoder = models.Sequential()
encoder.add(layers.Conv2D(32, (3, 3), activation='relu', padding = 'same', input_shape=(28, 28, 1)))
encoder.add(layers.MaxPooling2D((2, 2)))
encoder.add(layers.Conv2D(16, (3, 3), activation='relu', padding = 'same'))
encoder.add(layers.MaxPooling2D((2, 2)))
encoder.add(layers.Conv2D(8, (3, 3), activation='relu', padding = 'same', name = 'convOutput'))
encoder.add(layers.Flatten())
encoder.add(layers.Dense(encoding_dim, activation='relu'))

# 2
conv_output_shape = encoder.get_layer('convOutput').output_shape[1:]
dense_output_shape = conv_output_shape[0]*conv_output_shape[1]*conv_output_shape[2]

decoder = models.Sequential()
decoder.add(layers.Dense(dense_output_shape, input_shape=(encoding_dim,)))
decoder.add(layers.Reshape(conv_output_shape))
decoder.add(layers.Conv2D(8, (3, 3), activation='relu', padding = 'same'))
decoder.add(layers.UpSampling2D((2, 2)))
decoder.add(layers.Conv2D(16, (3, 3), activation='relu', padding = 'same'))
decoder.add(layers.UpSampling2D((2, 2)))
decoder.add(layers.Conv2D(32, (3, 3), activation='relu', padding = 'same'))
decoder.add(layers.Conv2D(1, (3, 3), activation='sigmoid', padding = 'same'))

# 3
autoencoder = models.Sequential()
autoencoder.add(encoder)
autoencoder.add(decoder)
```

```
In [ ]: encoder.summary()
decoder.summary()
autoencoder.summary()
```


Model: "sequential_30"

Layer (type)	Output Shape	Param #
conv2d_46 (Conv2D)	(None, 28, 28, 32)	320
max_pooling2d_22 (MaxPooling2D)	(None, 14, 14, 32)	0
conv2d_47 (Conv2D)	(None, 14, 14, 16)	4624
max_pooling2d_23 (MaxPooling2D)	(None, 7, 7, 16)	0
convOutput (Conv2D)	(None, 7, 7, 8)	1160
flatten_16 (Flatten)	(None, 392)	0
dense_47 (Dense)	(None, 32)	12576

=====
Total params: 18680 (72.97 KB)
Trainable params: 18680 (72.97 KB)
Non-trainable params: 0 (0.00 Byte)

Model: "sequential_31"

Layer (type)	Output Shape	Param #
dense_48 (Dense)	(None, 392)	12936
reshape_6 (Reshape)	(None, 7, 7, 8)	0
conv2d_48 (Conv2D)	(None, 7, 7, 8)	584
up_sampling2d_12 (UpSampling2D)	(None, 14, 14, 8)	0
conv2d_49 (Conv2D)	(None, 14, 14, 16)	1168
up_sampling2d_13 (UpSampling2D)	(None, 28, 28, 16)	0
conv2d_50 (Conv2D)	(None, 28, 28, 32)	4640
conv2d_51 (Conv2D)	(None, 28, 28, 1)	289

=====
Total params: 19617 (76.63 KB)
Trainable params: 19617 (76.63 KB)
Non-trainable params: 0 (0.00 Byte)

Model: "sequential_32"

Layer (type)	Output Shape	Param #
=====		

sequential_30 (Sequential)	(None, 32)	18680
sequential_31 (Sequential)	(None, 28, 28, 1)	19617

```
=====
Total params: 38297 (149.60 KB)
Trainable params: 38297 (149.60 KB)
Non-trainable params: 0 (0.00 Byte)
```

Task 3: Create the neural network model

Fit the model to the training data using the following hyper-parameters:

- adam optimizer
- binary_crossentropy loss function
- 20 training epochs
- batch size as 256
- set shuffle as True

Compile the model and fit ...

```
In [ ]: autoencoder.compile(optimizer='adam', loss='binary_crossentropy')
        history = autoencoder.fit(train_images_noisy, train_images_nor,
                                epochs=20,
                                batch_size=256,
                                shuffle=True)
```

```

Epoch 1/20
235/235 [=====] - 5s 11ms/step - loss: 0.2400
Epoch 2/20
235/235 [=====] - 3s 12ms/step - loss: 0.1413
Epoch 3/20
235/235 [=====] - 3s 11ms/step - loss: 0.1253
Epoch 4/20
235/235 [=====] - 3s 11ms/step - loss: 0.1179
Epoch 5/20
235/235 [=====] - 3s 11ms/step - loss: 0.1137
Epoch 6/20
235/235 [=====] - 3s 11ms/step - loss: 0.1106
Epoch 7/20
235/235 [=====] - 3s 12ms/step - loss: 0.1083
Epoch 8/20
235/235 [=====] - 3s 11ms/step - loss: 0.1065
Epoch 9/20
235/235 [=====] - 3s 11ms/step - loss: 0.1051
Epoch 10/20
235/235 [=====] - 3s 11ms/step - loss: 0.1038
Epoch 11/20
235/235 [=====] - 3s 11ms/step - loss: 0.1029
Epoch 12/20
235/235 [=====] - 3s 12ms/step - loss: 0.1021
Epoch 13/20
235/235 [=====] - 3s 11ms/step - loss: 0.1014
Epoch 14/20
235/235 [=====] - 3s 11ms/step - loss: 0.1008
Epoch 15/20
235/235 [=====] - 3s 11ms/step - loss: 0.1001
Epoch 16/20
235/235 [=====] - 3s 11ms/step - loss: 0.0998
Epoch 17/20
235/235 [=====] - 3s 11ms/step - loss: 0.0992
Epoch 18/20
235/235 [=====] - 2s 11ms/step - loss: 0.0989
Epoch 19/20
235/235 [=====] - 3s 11ms/step - loss: 0.0985
Epoch 20/20
235/235 [=====] - 2s 11ms/step - loss: 0.0981

```

Task 4: Create the neural network model (No need to write code, just run the following commands)

```

In [ ]: def showImages(input_imgs, encoded_imgs, output_imgs, size=1.5, groundTruth=None):

    numCols = 3 if groundTruth is None else 4

    num_images = input_imgs.shape[0]

    encoded_imgs = encoded_imgs.reshape((num_images, 1, -1))

    plt.figure(figsize=((numCols+encoded_imgs.shape[2])/input_imgs.shape[2])*size, num

```

```
pltIdx = 0
col = 0
for i in range(0, num_images):

    col += 1
    # plot input image
    pltIdx += 1
    ax = plt.subplot(num_images, numCols, pltIdx)
    plt.imshow(input_imgs[i].reshape(28, 28))
    plt.gray()
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)
    if col == 1:
        plt.title('Input Image')

    # plot encoding
    pltIdx += 1
    ax = plt.subplot(num_images, numCols, pltIdx)
    plt.imshow(encoded_imgs[i])
    plt.gray()
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)
    if col == 1:
        plt.title('Encoded Image')

    # plot reconstructed image
    pltIdx += 1
    ax = plt.subplot(num_images, numCols, pltIdx)
    plt.imshow(output_imgs[i].reshape(28, 28))
    plt.gray()
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)
    if col == 1:
        plt.title('Reconstructed Image')

    if numCols == 4:
        # plot ground truth image
        pltIdx += 1
        ax = plt.subplot(num_images, numCols, pltIdx)
        plt.imshow(groundTruth[i].reshape(28, 28))
        plt.gray()
        ax.get_xaxis().set_visible(False)
        ax.get_yaxis().set_visible(False)

        if col == 1:
            plt.title('Ground Truth')

plt.show()
```

```
In [ ]: num_images = 10

input_labels = test_labels[0:num_images]
I = np.argsort(input_labels)

input_imgs = test_images_noisy[I]
```

```
encoded_imgs = encoder.predict(test_images_noisy[1])
output_imgs = decoder.predict(encoded_imgs)

showImages(input_imgs, encoded_imgs, output_imgs, size=2, groundTruth=test_images_n
```

1/1 [=====] - 0s 63ms/step

1/1 [=====] - 0s 82ms/step

